

One shared visual-risk estimate, two hardware-feasible decisions


(a) Input




Tiny probe output



(b) Candidate risk · $\tau=0.066$




(c) E3 · TensorRT/CUDA · B=60 ms



Large output



(d) E1 · Hailo-8 · B=60 ms



Small output

