


(a) One shared elastic encoder-decoder



Three elastic axes

R_ℓ input resolution

w_ℓ active channel prefix (width)

d_ℓ active block prefix (depth)

weights shared by slicing;
batch-norm statistics kept per capacity

(b) Four static candidates extracted from the same weights

Tiny

384×192



$$w = 0.25 \cdot d = 2$$

0.033 M params · 0.206 GFLOPs

candidate-only p95 (ms)
AGX 1.03 · Hailo-8 4.01

Small

512×256



$$w = 0.50 \cdot d = 3$$

0.152 M params · 1.168 GFLOPs

candidate-only p95 (ms)
AGX 2.29 · Hailo-8 8.18

Medium

768×384



$$w = 0.75 \cdot d = 4$$

0.420 M params · 6.171 GFLOPs

candidate-only p95 (ms)
AGX 6.86 · Hailo-8 20.93

Large

1024×512



$$w = 1.00 \cdot d = 6$$

1.021 M params · 22.699 GFLOPs

candidate-only p95 (ms)
AGX 18.13 · Hailo-8 49.41

- encoder stage
- decoder stage
- 19-class head
- image / output map
- skip connection
- feature flow

each candidate

fixed-shape ONNX

TensorRT FP16 (AGX / NX / Orin)

Hailo-8 INT8 HEF